

# Database Design

## 1. Database Implementation

**\*FROM THIS STAGE WE RENAME TABLE “ILLNESSES” TO “CONDITIONS” FOR CONSISTENCY, SORRY FOR THE CONFUSION\***

In this stage we decide to implement our database on both local and GCP. We fixed the “Missed the Updates Relationship in the Schema.” problem in the previous stage by specifying the Update/Delete conditions. The following code is the DDL we used.

```
1 CREATE TABLE USER(  
2 user_id INT,  
3 sex VARCHAR(255),  
4 country VARCHAR(255),  
5 FirstName VARCHAR(255),  
6 LastName VARCHAR(255),  
7 password VARCHAR(255),  
8 age INT,  
9 PRIMARY KEY(user_id)  
10 );  
11  
12 CREATE TABLE CONDITIONS(  
13 trackable_id INT,  
14 name VARCHAR(255),  
15 description VARCHAR(255),  
16 treatment_options VARCHAR(255),  
17 amount_of_patient INT,  
18 average_age INT,  
19 PRIMARY KEY(trackable_id)  
20 );  
21  
22 CREATE TABLE SYMPTOMS(  
23 trackable_id INT,  
24 amount_of_patient INT,  
25 average_age INT,  
26 name VARCHAR(255),  
27 location VARCHAR(255),  
28 description VARCHAR(255),  
29 PRIMARY KEY(trackable_id)  
30 );  
31
```

```

32 CREATE TABLE REPORTING(
33 report_id INT,
34 report_type VARCHAR(255),
35 reported_symptom VARCHAR(255),
36 reported_description VARCHAR(255),
37 PRIMARY KEY(report_id)
38 );
39
40 CREATE TABLE WISHLIST(
41 wishlist_id INT,
42 illness_name VARCHAR(255),
43 illness_description VARCHAR(255),
44 user_id INT,
45 subscription_email VARCHAR(255),
46 PRIMARY KEY(wishlist_id),
47 FOREIGN KEY(user_id) REFERENCES USER(user_id) ON DELETE CASCADE ON UPDATE CASCADE
48 );
49
50 CREATE TABLE relate_to(
51 condition_id INT,
52 symptom_id INT,
53 condition_name VARCHAR(255),
54 symptom_name VARCHAR(255),
55 relation_count INT,
56 PRIMARY KEY(condition_id, symptom_id),
57 FOREIGN KEY(condition_id) REFERENCES CONDITIONS(trackable_id) ON DELETE CASCADE ON UPDATE CASCADE,
58 FOREIGN KEY(symptom_id) REFERENCES SYMPTOMS(trackable_id) ON DELETE CASCADE ON UPDATE CASCADE
59 );
60
61
62 CREATE TABLE diagnosed_with(
63 condition_id INT,
64 user_id int,
65 PRIMARY KEY(condition_id, user_id),
66 FOREIGN KEY(condition_id) REFERENCES CONDITIONS(trackable_id) ON DELETE CASCADE ON UPDATE CASCADE,
67 FOREIGN KEY(user_id) REFERENCES USER(user_id) ON DELETE CASCADE ON UPDATE CASCADE
68 );
69
70 CREATE TABLE suffer_from(
71 user_id INT,
72 symptom_id INT,
73 PRIMARY KEY(user_id, symptom_id),
74 FOREIGN KEY(user_id) REFERENCES USER(user_id) ON DELETE CASCADE ON UPDATE CASCADE,
75 FOREIGN KEY(symptom_id) REFERENCES SYMPTOMS(trackable_id) ON DELETE CASCADE ON UPDATE CASCADE
76 );

```

All of the tables in the database is shown here:

```

mysql> SHOW TABLES;
+-----+
| Tables_in_cs411 |
+-----+
| CONDITIONS      |
| REPORTING       |
| SYMPTOMS        |
| USER            |
| WISHLIST        |
| diagnosed_with  |
| relate_to       |
| suffer_from     |
+-----+
8 rows in set (0.00 sec)

```

Here is some deep look into the structure of three of our tables:

```
mysql> DESCRIBE SYMPTOMS;
```

Field	Type	Null	Key	Default	Extra
trackable_id	int(11)	NO	PRI	NULL	
amount_of_patient	int(11)	YES		NULL	
average_age	int(11)	YES		NULL	
name	varchar(255)	YES		NULL	
location	varchar(255)	YES		NULL	
description	varchar(255)	YES		NULL	

6 rows in set (0.00 sec)

```
mysql> DESCRIBE ILLNESES;
```

Field	Type	Null	Key	Default	Extra
trackable_id	int(11)	NO	PRI	NULL	
name	varchar(255)	YES		NULL	
description	varchar(255)	YES		NULL	
treatment_options	varchar(255)	YES		NULL	
amount_of_patient	int(11)	YES		NULL	
average_age	int(11)	YES		NULL	

6 rows in set (0.00 sec)

```
mysql> DESCRIBE relate_to;
```

Field	Type	Null	Key	Default	Extra
illness_id	int(11)	NO	PRI	NULL	
symptom_id	int(11)	NO	PRI	NULL	

2 rows in set (0.00 sec)

```
mysql> DESCRIBE CONDITIONS;
```

Field	Type	Null	Key	Default	Extra
trackable_id	int(11)	NO	PRI	NULL	
name	varchar(255)	YES		NULL	
description	varchar(255)	YES		NULL	
treatment_options	varchar(255)	YES		NULL	
amount_of_patient	int(11)	YES		NULL	
average_age	int(11)	YES		NULL	

6 rows in set (0.04 sec)

```
mysql> DESCRIBE relate_to;
```

Field	Type	Null	Key	Default	Extra
condition_id	int(11)	NO	PRI	NULL	
symptom_id	int(11)	NO	PRI	NULL	
condition_name	varchar(255)	YES		NULL	
symptom_name	varchar(255)	YES		NULL	
relation_count	int(11)	YES		NULL	

5 rows in set (0.00 sec)

For GCP it's all the same, here's a screenshot of the console:



## 2. Data Pre-Processing

The dataset we used is rich in information yet lacks organization, so we spent a great amount of time pre-processing it to make it usable. We took the following steps to organize the dataset into separate clean data.

### Conditions table and Symptoms table

Total number of unique conditions ('trackable\_name') tracked are 9443  
Total number of unique symptoms ('trackable\_name') tracked are 23157

Initially, we got 9443 unique conditions and 23157 symptoms in the dataset. To simplify the case, we choose the top 1000 conditions/symptoms by the number of patients reporting having those. Then we calculate the average age of those conditions/symptoms based on the original dataset.

Step 1: Cleaning up and splitting the dataset

```
1 import pandas as pd
2 import numpy as np
3 ## STEP 1: CLEAN UP AND SPLIT THE DATASET
4 res = pd.read_csv('export.csv')
5 df = pd.DataFrame(res)
6 df['user_id'] = pd.Categorical(df['user_id'])
7 df['user_id'] = df.user_id.cat.codes
8 df['age'] = df.age.replace(0.0, np.nan)
9
10 # Save data with trackable_type=="Symptom" or trackable_type=="Condition"
11 df_filtered_Symptom = df[df.trackable_type=="Symptom"]
12 df_filtered_Symptom.to_csv('Symptom.csv')
13 df_filtered_Condition = df[df.trackable_type=="Condition"]
14 df_filtered_Condition.to_csv('Condition.csv')
```

Step 2: Fetch top 1000 symptoms (same for conditions)

```

17 # # STEP 2: GET TOP 1000 SYMPTOMS
18 res = pd.read_csv('Symptom.csv')
19 # Create a DataFrame and filter for rows with "trackable_type" equal to "Symptom"
20 df = pd.DataFrame(res[res.trackable_type=="Symptom"])
21 # Get the top 10 frequently occurring symptoms
22 output = df.trackable_name.value_counts().head(1000)
23 # Calculate the mean age for each symptom in the top 10 list
24 mean_age = df[df.trackable_name.isin(output.index)].groupby("trackable_name").mean().sort_values(by="age")
25 # Print the mean age of each symptom in the top 1000 list, sorted by the top 1000 order
26 out = mean_age.loc[output.index]
27 out.to_csv('Sym1000age.csv')

```

Step 3: Count top 1000 conditions' number of patients (same for symptoms)

```

# STEP 4: GET TOP 1000 CONDITIONS COUNTS
res = pd.read_csv('Condition.csv')
df = pd.DataFrame(res)
print("Total number of unique Condition ('trackable_name') tracked are", df[df.trackable_type=="Condition"].trackable_name.nunique())
count = df[df.trackable_type=="Condition"].trackable_name.value_counts().head(1000)
count.to_csv('topCon.csv')

```

Step

4: Combine the above data and round up data to make things neat. **Conditions-Symptoms relation table**

For simplicity, we randomly chose 400000 data from the original dataset. Then we generated a table for the correlated relationship with the Symptoms and Conditions by counting how many patients reported having both simultaneously. We chose the correlations that have more than 500 records.

```

res = pd.read_csv('export.csv')
df = pd.DataFrame(res)
df = df.sample(n=400000)
# Create two DataFrames, one for Conditions and one for Symptoms
conditions = df[df['trackable_type'] == 'Condition'].rename(columns={'trackable_id': 'condition_id', 'trackable_name': 'condition_name'})
symptoms = df[df['trackable_type'] == 'Symptom'].rename(columns={'trackable_id': 'symptom_id', 'trackable_name': 'symptom_name'})
# Merge the two DataFrames on the user_id column
merged = pd.merge(conditions, symptoms, on='user_id')
# Group the merged DataFrame by user_id, condition_id, symptom_id, condition_name, and symptom_name, then count the occurrences
counts = merged.groupby(['user_id', 'condition_id', 'symptom_id', 'condition_name', 'symptom_name']).size().reset_index(name='count')
# Filter the counts DataFrame to only include rows where count is greater than or equal to 2
correlated = counts[counts['count'] >= 500]
print(correlated.shape[0])
correlated.to_csv('relate.csv')

```

### 3. Data Insertion

In this stage, we finished inserting data into table “CONDITIONS”, “SYMPTOMS” and “relate\_to”. The following shows the SQL commands we used for our local database. For the GCP, we simply click import.

```

1 LOAD DATA LOCAL INFILE "/Users/mac/Desktop/topCon.csv" INTO TABLE CS411.CONDITIONS
2 FIELDS TERMINATED BY ','
3 LINES TERMINATED BY '\n'
4 IGNORE 1 LINES;
5
6 LOAD DATA LOCAL INFILE "/Users/mac/Desktop/topSym.csv" INTO TABLE CS411.SYMPTOMS
7 FIELDS TERMINATED BY ','
8 LINES TERMINATED BY '\n'
9 IGNORE 1 LINES;
10
11 LOAD DATA LOCAL INFILE "/Users/mac/Desktop/relate.csv" INTO TABLE CS411.relate_to
12 FIELDS TERMINATED BY ','
13 LINES TERMINATED BY '\n'
14 IGNORE 1 LINES;

```

Here are some screenshots of our database:

Insert 1000 rows into TABLE “CONDITIONS”

```
mysql> SELECT COUNT(*) FROM CONDITIONS;
+-----+
| COUNT(*) |
+-----+
|      1000 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT * FROM CONDITIONS LIMIT 15;
```

trackable_id	name	description	treatment_options	amount_of_patient	average_age
2	ABHO			3768	31
7	Achalasia			223	48
8	Achilles tendinitis			179	41
10	Azoe			1831	38
19	Acute sinusitis			647	45
20	Addison's disease			316	33
21	Adenomyosis			473	26
26	Adult ADHD			1245	24
28	Agoraphobia			593	33
36	Allergies			8722	34
47	Anal fissure			153	34
50	Anaphylaxis			119	33
51	Anemia			1067	33
56	Ankylosing spondylitis			3026	38
57	Anorexia nervosa			648	24

```
15 rows in set (0.00 sec)
```

Insert 1000 rows into TABLE “SYMPTOMS”

```
mysql> SELECT COUNT(*) FROM SYMPTOMS;
+-----+
| COUNT(*) |
+-----+
|      1000 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> SELECT * FROM SYMPTOMS LIMIT 15;
```

trackable_id	amount_of_patient	average_age	name	location	description
1	33275	31	Abdominal pain		
3	1652	31	Allergy		
5	564	45	Anal itching		
6	510	34	Anemia		
8	61545	34	Anxiety		
9	985	37	Aphasia		
11	6806	34	Ankle pain		
12	37830	32	Back pain		
13	499	32	Bad breath		
16	4417	33	Arm pain		
17	835	36	Bladder spasms		
18	512	35	Bleeding gums		
22	625	30	Bloody nose		
25	6659	37	Blurred vision		
26	1263	44	Bowel incontinence		

```
15 rows in set (0.03 sec)
```

Insert 1000+ rows into TABLE “relate\_to”

```
mysql> SELECT COUNT(*) FROM relate_to;
+-----+
| COUNT(*) |
+-----+
|      1188 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM relate_to ORDER BY RAND() LIMIT 20;
+-----+-----+-----+-----+-----+
| condition_id | symptom_id | condition_name | symptom_name | relation_count |
+-----+-----+-----+-----+-----+
| 1199 | 242 | Acid Reflux | Fatigue | 5250 |
| 1186 | 279 | Mast Cell Activation Syndrome | Muscle spasms | 572 |
| 1174 | 121 | joint pain | Joint pain | 4389 |
| 1132 | 331 | POTS | muscle weakness | 520 |
| 421 | 13432 | Gastroparesis | being sick | 506 |
| 80 | 45 | Asthma | Constipation | 638 |
| 222 | 1890 | Chronic hives (urticaria) | Nose sores | 525 |
| 291 | 35 | Depression | Chest pain | 528 |
| 397 | 2762 | Fibromyalgia | hand and wrist pain | 504 |
| 6154 | 12247 | bowel not working very well | got shuff pouring out underneath | 675 |
| 532 | 54 | Idiopathic hypersomnia | Depression | 1089 |
| 291 | 515 | Depression | exhaustion | 528 |
| 1164 | 329 | Fatigue | impaired cognition | 1276 |
| 80 | 7004 | Asthma | GI Distress | 506 |
| 1201 | 674 | Chronic Pain | Neuropathy | 961 |
| 485 | 11505 | Hepatitis C | How Do You Feel Today? | 720 |
| 1493 | 8518 | ME/CFS | overexertion | 1421 |
| 269 | 243 | Crohn's disease | Stomach Pain | 532 |
| 1158 | 142 | OCD | Mood swings | 912 |
| 728 | 415 | Osteoarthritis | Achiness | 754 |
+-----+-----+-----+-----+-----+
20 rows in set (0.01 sec)
```

## 4. Advanced Queries

### QUERY1: Condition and symptom names that are related to each other

This query finds the average age of patients who have at least two symptoms in common, along with the names of the conditions and symptoms. This is useful because we want to find patients that have similar histories and then we can extrapolate more information based off of those similar histories. Perhaps even find a cause for the symptoms or derive risk factors from them.

#### The commands:

```
SELECT c.name AS condition_name, s.name AS symptom_name, AVG(s.average_age)
AS avg_age
```

```
FROM CONDITIONS c JOIN relate_to r1 ON c.trackable_id = r1.condition_id
```

```
JOIN SYMPTOMS s ON r1.symptom_id = s.trackable_id
```

```
JOIN relate_to r2 ON r1.symptom_id = r2.symptom_id AND r1.condition_id !=
r2.condition_id
```

```
GROUP BY c.name, s.name
```

```
HAVING COUNT(DISTINCT r2.condition_id) > 1;
```

## The results screenshot:

The screenshot shows a SQL IDE interface with a query editor and a results grid. The query editor contains the following SQL code:

```
1 • SELECT c.name AS condition_name, s.name AS symptom_name, AVG(s.average_age) AS a
2 FROM CONDITIONS c JOIN relate_to r1 ON c.trackable_id = r1.condition_id
3 JOIN SYMPTOMS s ON r1.symptom_id = s.trackable_id
4 JOIN relate_to r2 ON r1.symptom_id = r2.symptom_id AND r1.condition_id != r2.con
5 GROUP BY c.name, s.name
6 HAVING COUNT(DISTINCT r2.condition_id) > 1;
7
```

The results grid displays the following data:

condition_name	symptom_name	avg_age
Acid Reflux	NULL	NULL
Acid Reflux	Anhedonia	35.0000
Acid Reflux	Bloating	33.0000
Acid Reflux	Coccydynia	39.0000
Acid Reflux	Constipation	34.0000
Acid Reflux	Costochondritis	37.0000
Acid Reflux	Depression	33.0000
Acid Reflux	Diarrhea	34.0000
Acid Reflux	Fatigue	34.0000
Acid Reflux	Headache	33.0000
Acid Reflux	Hot flashes	40.0000
Acid Reflux	Hunger	37.0000
Acid Reflux	Hyperacusis	37.0000
Acid Reflux	Hypomania	36.0000
Acid Reflux	impaired cognit...	44.0000
Acid Reflux	Joint pain	33.0000
Acid Reflux	lightheadedness	32.0000
Acid Reflux	Migraine	33.0000
Acid Reflux	myalgia	44.0000
Acid Reflux	Nausea	32.0000
Acid Reflux	nerve pain	35.0000
Acid Reflux	Nosebleed	37.0000
Acid Reflux	Shortness of b...	34.0000
Acid Reflux	Stiffness	40.0000
Acid Reflux	Tachycardia	31.0000
Ankylosing spo...	Leg pain	32.0000
Anxiety	NULL	NULL

The interface includes a toolbar with icons for file operations, a 'Limit to 1000 rows' dropdown, and a sidebar with buttons for 'Result Grid', 'Form Editor', 'Field Types', 'Query Stats', and 'Execution Plan'. The status bar at the bottom indicates 'Result 1' and 'Read Only'.

## Indexing Analysis

Default query:



```

1 -> Filter: (count(distinct relate_to.condition_id) > 1) (actual time=21.133..26.635 rows=956 loops=1)
2   -> Group aggregate: count(distinct relate_to.condition_id), avg(symptoms.average_age) (actual
   time=21.132..26.566 rows=1039 loops=1)
3     -> Sort: c.name, s.name (actual time=21.117..21.801 rows=11736 loops=1)
4       -> Stream results (cost=1972.17 rows=4780) (actual time=0.041..13.004 rows=11736 loops=1)
5         -> Nested loop inner join (cost=1972.17 rows=4780) (actual time=0.039..10.026 rows=11736
           loops=1)
6           -> Nested loop inner join (cost=1141.35 rows=1193) (actual time=0.032..2.646 rows=1193
              loops=1)
7             -> Nested loop inner join (cost=723.80 rows=1193) (actual time=0.028..1.872 rows=1193
                  loops=1)
8               -> Covering index scan on r1 using rs_idx (cost=127.30 rows=1193) (actual
                   time=0.020..0.311 rows=1193 loops=1)
9                 -> Single-row index lookup on c using PRIMARY (trackable_id=r1.condition_id) (
                     cost=0.40 rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
10                  -> Single-row index lookup on s using PRIMARY (trackable_id=r1.symptom_id) (cost=0.25
                       rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
11                   -> Filter: (r1.condition_id <> r2.condition_id) (cost=0.25 rows=4) (actual
                       time=0.002..0.006 rows=10 loops=1193)
12                     -> Covering index lookup on r2 using rs_idx (symptom_id=r1.symptom_id) (cost=0.25
                           rows=4) (actual time=0.002..0.005 rows=11 loops=1193)
13
14

```

So this is the EXPLAIN ANALYZE result of the original query.

*Index 1:* CREATE INDEX symptomname\_idx on symptoms(name);

```

1 ▾ -> Filter: (count(distinct relate_to.condition_id) > 1) (actual time=22.808..28.489 rows=956 loops=1)
2 ▾   -> Group aggregate: count(distinct relate_to.condition_id), avg(symptoms.average_age) (actual
   time=22.806..28.416 rows=1039 loops=1)
3 ▾     -> Sort: c.name, s.name (actual time=22.791..23.526 rows=11736 loops=1)
4 ▾       -> Stream results (cost=1972.17 rows=4780) (actual time=0.069..14.320 rows=11736 loops=1)
5 ▾         -> Nested loop inner join (cost=1972.17 rows=4780) (actual time=0.065..11.225 rows=11736
           loops=1)
6 ▾           -> Nested loop inner join (cost=1141.35 rows=1193) (actual time=0.057..3.207 rows=1193
              loops=1)
7 ▾             -> Nested loop inner join (cost=723.80 rows=1193) (actual time=0.053..2.089 rows=1193
                  loops=1)
8               -> Covering index scan on r1 using rs_idx (cost=127.30 rows=1193) (actual
                   time=0.037..0.330 rows=1193 loops=1)
9                 -> Single-row index lookup on c using PRIMARY (trackable_id=r1.condition_id) (
                     cost=0.40 rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
10                  -> Single-row index lookup on s using PRIMARY (trackable_id=r1.symptom_id) (cost=0.25
                       rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
11 ▾                   -> Filter: (r1.condition_id <> r2.condition_id) (cost=0.25 rows=4) (actual
                       time=0.002..0.006 rows=10 loops=1193)
12                     -> Covering index lookup on r2 using rs_idx (symptom_id=r1.symptom_id) (cost=0.25
                           rows=4) (actual time=0.002..0.005 rows=11 loops=1193)
13
14

```

To begin we make a simple index on the symptom name to see if it makes any difference. You can see that the times actually increase, but the costs themselves are the exact same as before. So this is not actually improving the efficiency of the query but seems to be making the times worse anyway.

*Index 2:* CREATE INDEX cidx on relate\_to(condition\_id, symptom\_id)

Let's try making an index for the relate\_to table instead. We will also try using two columns this time, condition\_id and symptom\_id

```

1 ▾ -> Filter: (count(distinct relate_to.condition_id) > 1) (actual time=20.481..25.758 rows=956 loops=1)
2 ▾   -> Group aggregate: count(distinct relate_to.condition_id), avg(symptoms.average_age) (actual
    time=20.479..25.686 rows=1039 loops=1)
3 ▾   -> Sort: c.name, s.name (actual time=20.463..21.170 rows=11736 loops=1)
4 ▾     -> Stream results (cost=1966.17 rows=4780) (actual time=0.079..12.541 rows=11736 loops=1)
5 ▾       -> Nested loop inner join (cost=1966.17 rows=4780) (actual time=0.075..9.661 rows=11736 loops=1)
6 ▾         -> Nested loop inner join (cost=1135.35 rows=1193) (actual time=0.067..2.621 rows=1193
            loops=1)
7 ▾           -> Nested loop inner join (cost=717.80 rows=1193) (actual time=0.049..1.913 rows=1193
                loops=1)
8 ▾             -> Covering index scan on r1 using rs_idx (cost=121.30 rows=1193) (actual
                time=0.036..0.325 rows=1193 loops=1)
9 ▾               -> Single-row index lookup on c using PRIMARY (trackable_id=r1.condition_id) (
                cost=0.40 rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
10 ▾                -> Single-row index lookup on s using PRIMARY (trackable_id=r1.symptom_id) (cost=0.25
                    rows=1) (actual time=0.000..0.000 rows=1 loops=1193)
11 ▾                 -> Filter: (r1.condition_id <> r2.condition_id) (cost=0.25 rows=4) (actual
                    time=0.002..0.005 rows=10 loops=1193)
12 ▾                   -> Covering index lookup on r2 using rs_idx (symptom_id=r1.symptom_id) (cost=0.25
                        rows=4) (actual time=0.002..0.005 rows=11 loops=1193)
13

```

So immediately there appears to be a boost in overall time (from 21.1 seconds to 20.4), and also we see the costs have lowered (from 1972 to 1966 for the top level operation. Some of the other operations have also decreased (1141 to 1135 for the nested inner loop join), 723 to 717, etc. lots of slight improvements). The structure of the query still hasn't changed though, there are still 12 operations total.

*Index 3: CREATE INDEX cidx on c(condition\_id, symptom\_id)*

```

1 ▾ -> Filter: (count(distinct relate_to.condition_id) > 1) (actual time=21.018..26.373 rows=956 loops=1)
2 ▾   -> Group aggregate: count(distinct relate_to.condition_id), avg(symptoms.average_age) (actual
    time=21.016..26.303 rows=1039 loops=1)
3 ▾   -> Sort: c.name, s.name (actual time=20.997..21.679 rows=11736 loops=1)
4 ▾     -> Stream results (cost=1787.22 rows=4780) (actual time=0.045..12.581 rows=11736 loops=1)
5 ▾       -> Nested loop inner join (cost=1787.22 rows=4780) (actual time=0.043..9.651 rows=11736 loops=1)
6 ▾         -> Nested loop inner join (cost=956.40 rows=1193) (actual time=0.033..2.594 rows=1193
            loops=1)
7 ▾           -> Nested loop inner join (cost=538.85 rows=1193) (actual time=0.029..1.907 rows=1193
                loops=1)
8 ▾             -> Covering index scan on r1 using rs_idx (cost=121.30 rows=1193) (actual
                time=0.021..0.309 rows=1193 loops=1)
9 ▾               -> Single-row index lookup on c using PRIMARY (trackable_id=r1.condition_id) (
                cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=1193)
10 ▾                -> Single-row index lookup on s using PRIMARY (trackable_id=r1.symptom_id) (cost=0.25
                    rows=1) (actual time=0.000..0.000 rows=1 loops=1193)
11 ▾                 -> Filter: (r1.condition_id <> r2.condition_id) (cost=0.25 rows=4) (actual
                    time=0.002..0.005 rows=10 loops=1193)
12 ▾                   -> Covering index lookup on r2 using rs_idx (symptom_id=r1.symptom_id) (cost=0.25
                        rows=4) (actual time=0.002..0.005 rows=11 loops=1193)
13

```

With this last attempt, the time actually goes up again, but the costs are significantly lower (by almost 200 points). You can see that the first “Stream results” operation goes from 1966 to 1787, which is better than the previous

two indexes and also better than the original query as well. Similarly, all the other operations have much lower costs (except for the last few which seem to be constant at a cost of 0.25). But overall if we had to choose one query it would be this one since it is the one showing the biggest improvements from the original.

### **QUERY 2: Select symptom that occurs in the most conditions**

This query retrieves the possible symptoms in the database, ordered by the frequency of their occurrence in conjunction with the conditions. Since the relate\_to table, brings the symptom and condition data together, we utilized it to find how many instances of the symptoms occur across the various conditions using COUNT.

```

1 SELECT r.symptom_id , COUNT(r.symptom_id) AS cnt
2 FROM relate_to r
3 GROUP BY r.symptom_id
4 ORDER BY COUNT(r.symptom_id) DESC;

```

CLOUD SHELL Terminal (cs411-380104) × +

symptom_id	cnt
98	39
152	38
242	34
45	25
63	23
8	22
121	22
145	21
275	21
54	19
12	17
35	17
56	14
399	14
257	13
515	13
191	12
329	11
2605	11
84	10
265	10
681	10

## Indexing Analysis

Default first:

```

# EXPLAIN
-> Sort: cnt DESC (actual time=1.262..1.269 rows=268 loops=1)
  -> Stream results (cost=240.60 rows=268) (actual time=0.400..0.754 rows=268 loops=1)
    -> Group aggregate: count(r.symptom_id) (cost=240.60 rows=268) (actual time=0.395..0.721 rows=268 loops=1)
      -> Covering index scan on r using symptom_id (cost=121.30 rows=1193) (actual time=0.260..0.526 rows=1193 loops=1)

```

So about 1.2 seconds

*Index 1:* CREATE INDEX rs\_idx on relate\_to(symptom\_id);

Let's try indexing symptom\_id:

```
# EXPLAIN
-> Sort: cnt DESC (actual time=0.518..0.527 rows=268 loops=1)
  -> Stream results (cost=240.60 rows=268) (actual time=0.106..0.448 rows=268 loops=1)
    -> Group aggregate: count(r.symptom_id) (cost=240.60 rows=268) (actual time=0.103..0.418 rows=268 loops=1)
      -> Covering index scan on r using rs_idx (cost=121.30 rows=1193) (actual time=0.097..0.349 rows=1193 loops=1)
```

Okay, so the time dropped about 0.7 seconds. Quite good.

*Index 2:* CREATE INDEX rc\_idx on relate\_to(condition\_id);

Well, there's not much else to index since in the query above we only really checked one field. So we will just try other fields in the relate\_to table and see if changing them gives any difference.

```
# EXPLAIN
-> Sort: cnt DESC (actual time=0.518..0.527 rows=268 loops=1)
  -> Stream results (cost=240.60 rows=268) (actual time=0.106..0.448 rows=268 loops=1)
    -> Group aggregate: count(r.symptom_id) (cost=240.60 rows=268) (actual time=0.103..0.418 rows=268 loops=1)
      -> Covering index scan on r using rs_idx (cost=121.30 rows=1193) (actual time=0.097..0.349 rows=1193 loops=1)
```

Strangely enough, it seems to still speed up the queries.

*Index 3:* CREATE INDEX rsn\_idx on relate\_to(symptom\_name);

```
# EXPLAIN
-> Sort: cnt DESC (actual time=0.629..0.636 rows=268 loops=1)
  -> Stream results (cost=240.60 rows=268) (actual time=0.055..0.537 rows=268 loops=1)
    -> Group aggregate: count(r.symptom_id) (cost=240.60 rows=268) (actual time=0.053..0.490 rows=268 loops=1)
      -> Covering index scan on r using rs_idx (cost=121.30 rows=1193) (actual time=0.046..0.389 rows=1193 loops=1)
```

And for the final one, it's a bit slower at 0.629 sec. The costs for all three are the same, so there's not much to comment on there. The query itself is perhaps a little too simple to benefit much from indexing. Or perhaps, we should try combinations of columns from different tables. But with the current strategies, we don't seem to get much improvement.